

BALAJI BERI

+91 7032581740 — balajiberi13@gmail.com — LinkedIn — LeetCode — Portfolio

Data Engineer with 2.8+ years of experience building and maintaining **ETL/ELT data pipelines**, optimizing **SQL queries**, and ensuring **data quality, monitoring, and troubleshooting**. Experienced in **Azure Data Factory (ADF)**, **Python**, **ADLS Gen2**, **Databricks (PySpark)**, and data warehousing concepts.

SKILLS

Languages: SQL, Python, PySpark

Cloud & Tools: Azure Data Factory (ADF), ADLS Gen2, Azure SQL, Databricks, Airflow, Azure DevOps, Azure Functions, Azure Key Vault, Azure Synapse, Azure Logic Apps

Data Engineering: ETL/ELT Pipelines, Data Modeling, Incremental Loads, CI/CD, Git, Stored Procedures

Core Concepts: Data Quality, Pipeline Monitoring, Performance Optimization, collaboration, Agile methodology, Logging, Troubleshooting, Data Structures, Data Governance

EDUCATION

B.Tech in Electronics and Communication Engineering

Sri Venkateswara University, Tirupati

2020 - 2024

EXPERIENCE

Quaxigma — Analyst (Data & AI)

Dec 2023 – Present

- Designed and developed scalable data ingestion pipelines to process large volumes of enterprise data across **7+ ERP systems** into ADLS Gen2.
- Engineered high-throughput ingestion workflows processing **2.5 TB** of Dailpad data using **AsyncIO and streaming** strategies for resilient and fault-tolerant data extraction.
- Built and implemented optimized **data models and ETL workflows** for manufacturing analytics, enabling efficient reporting across finance and operations.
- Implemented **data quality checks, monitoring, and validation frameworks** using SQL and Python, reducing data discrepancies by **25%** and eliminating **90%** manual validations.
- Assisted in **troubleshooting data pipeline failures and performance issues**, improving pipeline reliability.
- Built scalable **API-driven ETL pipelines** to ingest data from multiple ERP endpoints into ADLS Gen2 using memory-efficient streaming ingestion, reducing memory utilization and ingestion time by **50%**.
- Modernized legacy SQL Warehouse stored procedures by migrating transformations to **Databricks (PySpark)**, leveraging Spark optimization techniques such as **partitioning and caching** to reduce processing time by **40%**.

PROJECTS

Bing News Analytics — Microsoft Fabric, Spark

GitHub

- Built a scalable end-to-end data pipeline on **Microsoft Fabric** using Data Factory, **Synapse** (Spark), and Lakehouse, processing 1M+ news records/day with incremental ingestion, reducing data latency by 60% compared to full reloads.
- Applied sentiment analysis using Synapse **Data Science** and text analytics **ML** models to derive insights from unstructured news data.
- Delivered production-grade analytics via Power BI dashboards, pipeline orchestration, **monitoring**, and alerts, cutting pipeline failures by 30% and optimizing compute usage to achieve 25% cost savings.

ERP Lakehouse Analytics Platform — Azure Data Factory, Databricks, ADLS Gen2

GitHub

- Built scalable **API-driven ingestion pipelines** using **Azure Functions** and **Python** to extract and process Business Central ERP data into ADLS Gen2, improving ingestion scalability.
- Secured secrets, API credentials, and connection strings using **Azure Key Vault** integrated with **Azure Data Factory** and **Databricks**, enabling centralized credential management and improving pipeline security compliance.
- Designed and maintained **Bronze, Silver, and Gold (Medallion)** layer transformations in Databricks using **PySpark** and **Delta Lake** principles.
- Built fault-tolerant **ELT pipelines** leveraging **Delta Lake** and partition-based processing techniques for large-scale ERP workloads, improving pipeline performance and scalability for high-volume data processing.
- Integrated data from multiple Business Central API endpoints into a centralized **Lakehouse architecture** for downstream analytics and reporting, reducing data modeling efforts for BI teams by **70%**.

CERTIFICATIONS

Microsoft Fabric Data Engineer Associate (**DP-700**)

Microsoft Fabric Analytics Engineer Associate (**DP-600**)

Microsoft Azure Fundamentals (**AZ-900**)

SQL (Advanced) – HackerRank

ACHIEVEMENTS

Secured **2nd place** in District Level at the 26th National Mathematics Olympiad (2015)

Achieved **3rd rank among 200+ participants** in a Coding Contest at SVUCE (2022)